



KULeuven

Department of
Computer Science

ADVANCED CAPITA SELECTA
ARTIFICIAL INTELLIGENCE (H02A8A)
Contemporary AI topics
Report

Jeroen Craps (r0292642)
Jorik De Waen (r0303087)

Academic year 2016–2017

I. INTRODUCTION

February 1996, **Deep Blue**¹ wins the first ever game of Chess as a computer against the world champion at the time Garry Kasparov. March 2016, **AlphaGo** defeats 18-time world champion Lee Sedol in a five-game Go match with a score of 4-1.

Originally the biggest test for an artificial intelligence was the Turing test. The test requires that a human being is unable to distinguish the machine from another human being during a conversation where it has to answer to questions. A limitation of this test is that there is no objective way to measure the progress towards the goals of AI [3].

Computers have been able to perform tasks better than humans like path planning, finding patterns and playing games. After 20 years of progress, computers are now able to defeat the human mind at very complex games, but can it solve the same questions as asked in a 4th grade exam?

Even though 4th grade exams are trivial to solve for humans, they present an enormous challenge for our current AI systems. Researchers are actively working on this problem as this is seen as a key component of any new measurement of artificial intelligence [3].

There are several reasons behind this. It has all the requirements of a test [3]: Accessible, Comprehensible, Measurable and offer a graduated progression for simple everyday things to deeper understanding of subjects. Also to answer these questions a significant improvement in language understanding and the modelling of the world are required. In this report we will be discussing some of the current methods that are being used to improve the performance of AI on these kind of tests.

¹The Chess playing computer designed by IBM.

II. OVERVIEW

To learn the required knowledge artificial intelligence needs a source of data to work with. These data will almost exclusively consist of scientific papers and elementary school books. Currently two separate directions are being pursued, image recognition [1] and text based knowledge extraction [4, 5].

Both streams will be explained in this report, but the focus will be put on the latter. The first will try to retrieve images from scientific papers and correctly label it by using the caption that is included with the table or figure. In the second method the goal is to extract logical statements from the text. By combining these logical statements into a knowledge graph [2]. TODO: DIT IS NIET HELEMAAL CORRECT While answering questions the keywords of the text will be extracted and filled into the Knowledge graph. The multiple choice questions will be reformed to multiple true or false questions. Every answer will lead to a percentage which will define how true the statement was. The answer with the largest percentage is considered to be the best answer for this question.

III. IMAGE RETRIEVAL

In current academic documents figures and tables are key sources of information, i.e. taking a look at a table in a paper can quickly summarize the work that has been done. However, these are not currently used in academic search engines. To better answer any kind of question the use of figure and tables is encouraged.

When there is a focus on scientific papers, it has been shown that **captions** are the key elements to identify figures and tables. The difference between body text and captions has been proven to be relatively easy to detect [1]. The search is done by looking for keywords that are likely to start a caption. Afterwards false positives are removed by applying a filter to

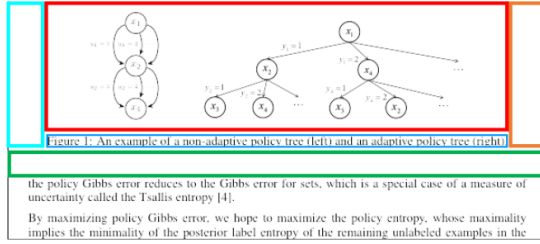


Figure 1: An example of separation into regions

them. This filter is focused on a particular format convention. This process of these filters is repeated until no false positives remain.

The second step is to classify every part of the paper as a certain region: caption, body text, graphical element or figure text. Caption regions are made from the captions and the following lines of text (if there are any). Poppler’s algorithm is used to define body text or figure text. Additional heuristics are applied for improved accuracy. To find the graphical elements, the pdf is directly parsed. Internally PDF’s make use of various “operators” that draw elements on the page. The algorithm uses the bounding boxes defined by the operators that the PDF would use to draw. The boxes can be merged if required to form the complete graphical region. An example of this can be seen in Figure 1.

To assign captions/titles to figures or tables clustering is utilized. These clusters would be pruned and additional rules for increased performance. Questions that would match up with elements in a caption can now be linked to a corresponding image which might contain useful information to answer the question.

IV. TEXT BASED KNOWLEDGE EXTRACTION

Texts from books and other sources often contain a large amount of information. Natural languages are flexible enough to communicate complex concepts, intricate relations and more. However, AI systems currently have a hard

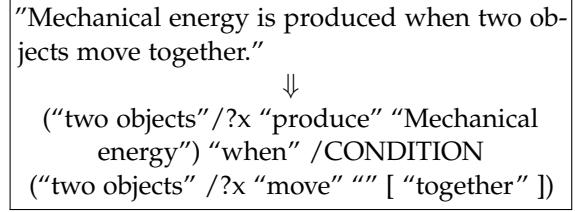


Figure 2: An example of an extraction

time accessing and reasoning with the knowledge contained in texts.

Even though interpreting natural is, so far, an unsolved problem, the limited scope of solving science tests brings a solution within reach. A significant part of the knowledge can be expressed with relations like “X causes Y”, “X is part of Y”, “X is an example of Y” or “X [verb] Y”. The algorithm by Clark et al. [2] uses a hand-crafted set of extraction rules to generate so-called extractions from the text. Figure 2 shows an example of this. Extractions are a semi-formal data structure, but they can easily be translated into formal representation. Many of these rules combined form a knowledge base.

V. ANSWERING QUESTIONS

Answering a question is a two-fold problem: The first part is actually understanding the question. A question is a query for a piece of knowledge, so the text of the question needs to be translated to a formal language. The second part is actually answering the question. Once the query is expressed in the right format, it needs to be executed on the knowledge base. Not every answer is stated explicitly in the knowledge base, so logical inference is required to combine multiple rules into a single answer.

i. Lexical analysis of questions

The approach chosen by Krishnamurthy [4] to parse the question is to actually reconstruct it from logical rules. The model is trained with

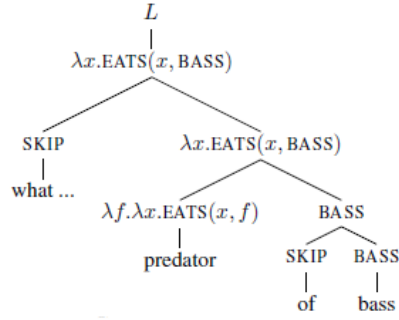


Figure 3: The parse tree for the question “What is the predator of bass?” During training, the algorithm works from the top down to build the lexical rules. When the model is actually applied, these lexical rules are applied from the bottom up to come up with the final interpretation.

manually translated questions. This allows the model to come up with possible rules and the likelihood that words and combinations of words translate into those rules. When the model is executed with a new question, it can look up the words and combine their corresponding rules from the bottom up. The probabilities calculated during training are used to decide at each step which of the options is the most likely.

ii. Querying the knowledge graph

Once the meaning of the question is clear, an answer needs to be found in the knowledge base. The algorithm by Li and Clark [5] is one way to achieve that goal. After parsing the question, they extract keywords. Background knowledge is added, which results in a knowledge graph of concepts connected by the relations found in the knowledge base. These concepts are the keywords themselves, but also other related words which provide context. The questions are multiple choice, so each option is placed within that graph and connected to the other concepts. The answer that

“Animals get energy for growth and repair from (A) food (B) ...

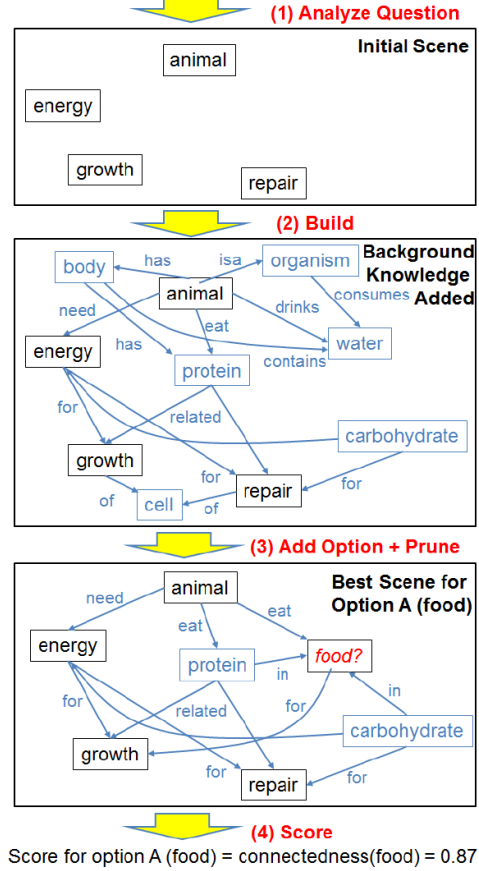


Figure 4: An example of a question being answered using a knowledge graph.

is the most connected with the concepts in the knowledge graph is chosen as the answer.

VI. CONCLUSION

To conclude this report, we can say that the current scientific progress is hopeful and improving rapidly. Texts can be interpreted to reasonable extent and converted to logical rules. Images can be extracted with the corresponding title and caption. With every iteration the scores that the artificial intelligence is able to achieve increases.

REFERENCES

- [1] Christopher Clark and Santosh Divvala. Pdffigures 2.0: Mining figures from research papers. In *Proceedings of the 16th ACM/IEEE-CS on Joint Conference on Digital Libraries*, JCDL '16, pages 143–152, New York, NY, USA, 2016. ACM.
- [2] Peter Clark, Niranjana Balasubramanian, Sumithra Bhakthavatsalam, Kevin Humphreys, and Kinlead. Automatic construction of inference-supporting knowledge bases. In *4th Workshop on Automated Knowledge Base Construction (AKBC)*, 2014.
- [3] Peter Clark and Oren Etzioni. My computer is an honor student - but how intelligent is it? standardized tests as a measure of ai. *AI Magazine*, 2016.
- [4] Jayant Krishnamurthy. Probabilistic models for learning a semantic parser lexicon. In *Proceedings of NAACL-HLT*, pages 606–616, 2016.
- [5] Yang Li and Peter Clark. Answering elementary science questions by constructing coherent scenes using background knowledge. In *EMNLP*, pages 2007–2012, 2015.